

# python-csp: bringing OCCAM to Python

Sarah Mount

Europython 2010

## 1 Ancient history

## 2 python-csp: features and idioms

- Process creation
- Parallel, sequential and repeated processes
- Communication via channel read / writes
- More channels: ALTing and choice
- More channels: poisoning
- Producer / consumer or worker / farmer models


## 3 Using built-in processes

## 4 Future challenges



## This talk...


is all about the python-csp project. There is a similar project called PyCSP, these will merge over the next few months. Current code base is here: <http://code.google.com/p/python-csp>  
Please do contribute bug fixes / issues / code / docs / rants / ...



## INMOS B0042 Board © David May

The Transputer was the original “multicore” machine and was built to run the OCCAM language – a realisation of CSP.

<http://www.cs.bris.ac.uk/~dave/transputer.html>



## Tony Hoare

Invented CSP, OCCAM, Quicksort and other baddass stuff. You may remember him from such talks as his Europython 2009 Keynote.

# OCCAM

OCCAM lives on in it's current implementation as OCCAM- $\pi$


<http://www.cs.kent.ac.uk/projects/ofa/kroc/>

It can run on Lego bricks, Arduino boards and other things. Like Python it uses indentation to demark scope. Unlike Python OCCAM is MOSTLY IN UPPERCASE.


Ancient history  
python-csp: features and idioms  
Using built-in processes  
Future challenges

## CSP and message passing

This next bit will be revision for many of you!



## Most of us think of programs like this



<http://xkcd.com/205/>



## This is the old “standard” model!

- Multicore will soon take over the world (if it hasn't already).
- Shared memory concurrency / parallelism is hard:
  - Race hazards
  - Deadlock
  - Livelock
  - Tracking which lock goes with which data and what order locks should be used.
  - Operating Systems are old news!


# Message passing

## Definition

In *message passing* concurrency (or parallelism) “everything” is a process and no data is shared between processes. Instead, data is “passed” between *channels* which connect processes together. Think: OS processes and UNIX pipes.


- The *actor* model (similar to Kamaelia) uses asynchronous channels. Processes write data to a channel and continue execution.
- CSP (and  $\pi$ -calculus) use synchronous channels, or *rendezvous*, where a process writing to a channel has to wait until the value has been read by another process before proceeding.

## Process diagram



## Why synchronous channels?

Fig. 1.



Leslie Lamport (1978). "Time, clocks, and the ordering of events in a distributed system". Communications of the ACM 21 (7)

# Hello World!

```
1 from csp.cspprocess import *
2
3 @process
4 def hello():
5     print 'Hello CSP world!'
6
7 hello().start()
```

## Example of process creation: web server

```
1 @process
2 def server(host, port):
3     sock = socket.socket(...)
4     # Set up sockets
5     while True:
6         conn, addr = sock.accept()
7         request = conn.recv(4096).strip()
8         if request.startswith('GET'):
9             ok(request, conn).start()
10        else:
11            not_found(request, conn).start()
```

## Example of process creation: web server

```
1 @process
2 def not_found(request, conn):
3     page = '<h1>File not found</h1>'
4     conn.send(response(404,
5                        'Not Found',
6                        page))
7     conn.shutdown(socket.SHUT_RDWR)
8     conn.close()
9     return
```

## Combining processes: in parallel

```
1      # With a Par object:
2      Par(p1, p2, p3, ... pn).start()
3      # With syntactic sugar:
4      p1 // (p2, p3, ... pn)
5      # RHS must be a sequence type.
```



## Combining processes: repeating a process sequentially

```
1 @process
2 def hello():
3     print 'Hello CSP world!'
4 if __name__ == '__main__':
5     hello() * 3
6     2 * hello()
```

# Channels

```
1 @process
2 def client(chan):
3     print chan.read()
4 @process
5 def server(chan):
6     chan.write('Hello CSP world!')
7
8 if __name__ == '__main__':
9     ch = Channel()
10    Par(client(ch), server(ch)).start()
```

# What you need to know about channels

- Channels are currently UNIX pipes
- This will likely change
- Channel rendezvous is “slow” compared with JCSP:  
 $200\mu s$  vs  $16\mu s$
- This will be fixed by having channels written in C
- Because channels are, at the OS level, open “files”, there can only be a limited number of them (around 300 on my machine)
- This makes me sad :-)

## Message passing an responsiveness

- This isn't about speed.
- Sometimes, you have several channels and reading from them all round-robin may reduce responsiveness if one channel read blocks for a long time.

This is BAD:

```
1 @process
2 def foo(channels):
3     for chan in channels:
4         print chan.read()
```

# ALTing and choice

## Definition

ALTing, or non-deterministic choice, *selects* a channel to read from from those channels which are ready to read from.

You can do this with (only) two channels:

```
1 @process
2 def foo(c1, c2):
3     print c1 | c2
4     print c1 | c2
```

## ALTING proper (many channels)

```
1 @process
2 def foo(c1, c2, c3, c4, c5):
3     alt = Alt(c1, c2, c3, c4, c5)
4     # Choose random available offer
5     print alt.select()
6     # Avoid last selected channel
7     print alt.fair_select()
8     # Choose channel with lowest index
9     print alt.priselect()
```

## Syntactic sugar for ALTing

```
1 @process
2 def foo(c1, c2, c3, c4, c5):
3     gen = Alt(c1, c2, c3, c4, c5)
4     print gen.next()
5     print gen.next()
6     print gen.next()
7     print gen.next()
8     print gen.next()
```

## ALing: when you really, really must return right now

```
1 @process
2 def foo(c1, c2, c3):
3     # Skip() will /always/ complete
4     # a read immediately
5     gen = Alt(c1, c2, c3, Skip())
6     print gen.next()
7     ...
```



# Channel poisoning

## Definition


Idiomatically in this style of programming most processes contain infinite loops. *Poisoning* a channel asks all processes connected to that channel to shut down automatically. Each process which has a poisoned channel will automatically poison any other channels it is connected to. This way, a whole graph of connected processes can be terminated, avoiding race conditions.

```
1 @process
2 def foo(channel):
3     ...
4     channel.poison()
```

Ancient history  
**python-csp: features and idioms**  
Using built-in processes  
Future challenges

Process creation  
Parallel, sequential and repeated processes  
Communication via channel read / writes  
More channels: ALTING and choice  
More channels: poisoning  
Producer / consumer or worker / farmer models

## A bigger example: Mandelbrot generator




```
1 @process
2 def main(IMSIZE):
3     channels = []
4     # One producer + channel for each image column
5     for x in range(IMSIZE[0]):
6         channels.append(Channel())
7         processes.append(mandelbrot(x, ... channel
8     processes.insert(0, consume(IMSIZE, channels))
9     mandel = Par(*processes)
10    mandel.start()
```

```
1 @process
2 def mandelbrot(xcoord, cout):
3     # Do some maths here
4     cout.write(xcoord, column_data)
```

```
1 @process
2 def consumer(cins):
3     # Initialise PyGame...
4     gen = len(cins) * Alt(*cins)
5     for i in range(len(cins)):
6         xcoord, column = gen.next()
7         # Blit that column to screen
8     for event in pygame.event.get():
9         if event.type == pygame.QUIT:
10             for channel in cins:
11                 channel.poison()
12             pygame.quit()
```

# Oscilloscope




```
1 @process
2 def Oscilloscope(inchan):
3     # Initialise PyGame
4     ydata = []
5     while not QUIT:
6         ydata.append(inchan.read())
7         ydata.pop(0)
8         # Blit data to screen...
9     # Deal with events
```

```
1 from csp.builtins import Sin
2 def trace_sin():
3     chans = Channel(), Channel()
4     Par(GenerateFloats(chans[0]),
5         Sin(chans[0], chans[1]),
6         Oscilloscope(chans[1])).start()
7     return
```



## Process diagram



```
1 def trace_mux():
2     chans = [Channel() for i in range(6)]
3     par = Par(GenerateFloats(chans[0]),
4               Delta2(chans[0], chans[1],
5                      chans[2]),
6               Cos(chans[1], chans[3]),
7               Sin(chans[2], chans[4]),
8               Mux2(chans[3], chans[4],
9                   chans[5]),
10                  Oscilloscope(chans[5]))
11     par.start()
```

## Wiring

